

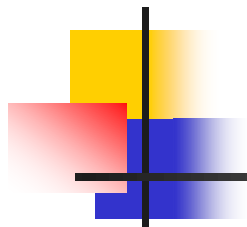
# Designing Web Applications

## HTTP

*Francesco Marcelloni*

Department of Information Engineering  
University of Pisa  
ITALY






# Acknowledgement

I acknowledge that some slides have been adapted from (or are just identical to) the slides provided with the following book:

Thee Computer Networking: A Top Down Approach ,  
5th edition. Jim Kurose, Keith Ross  
Addison-Wesley, April 2009.



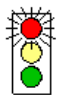








# TCP/IP Stack

From Computer Desktop Encyclopedia  
© 2003 The Computer Language Co. Inc.



OSI MODEL		TCP / IP
7	 <b>Application Layer</b> Type of communication: E-mail, file transfer, client/server.	FTP, Telnet, HTTP, SNMP, DNS, OSPF, RIP, Ping, Traceroute
6	 <b>Presentation Layer</b> Encryption, data conversion: ASCII to EBCDIC, BCD to binary, etc.	
5	 <b>Session Layer</b> Starts, stops session. Maintains order.	
4	 <b>Transport Layer</b> Ensures delivery of entire file or message.	TCP (delivery ensured)  UDP (delivery NOT ensured)
3	 <b>Network Layer</b> Routes data to different LANs and WANs based on network address.	IP  (ICMP, IGMP, ARP, RARP)
2	 <b>Data Link (MAC) Layer</b> Transmits packets from node to node based on station address.	
1	 <b>Physical Layer</b> Electrical signals and cabling.	



# HTTP Overview

## First some jargon


- Web page consists of objects
- Objects can be HTML files, JPEG images, Java applets, audio files,...
- Web page consists of base HTML-file which includes several referenced objects
- Each object is addressable by a URI
- Example URI:

`www.someschool.edu/someDept/pic.gif`

host name

path name





# HTTP Overview

`http://guest:secret@www.ietf.org:80/html.charters/wg-dir.html?sess=1#Applications_Area`


protocol	...	http
username	.....	guest
password	.....	secret
host	.....	www.ietf.org
port	.....	80
path	.....	/html.charters
file	.....	wg-dir.html
query	.....	sess=1
fragment	.....	Applications_Area



# HTTP Overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - *client*: browser that requests, receives, "displays" Web objects
  - *server*: Web server sends objects in response to requests





# HTTP Overview

## Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## HTTP is "stateless"

- server maintains no information about past client requests

Protocols that maintain "state" are complex! aside

- past history (state) must be maintained
- if server/client crashes, their views of "state" may be inconsistent, must be reconciled

# HTTP Connections

## Nonpersistent HTTP

- At most one object is sent over a TCP connection.

## Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server.



# Nonpersistent HTTP

Suppose user enters URL

(contains text, references to 10 jpeg images)

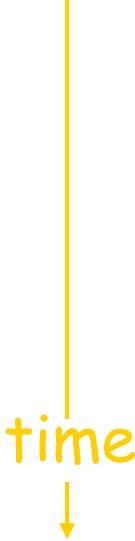
`www.someSchool.edu/someDepartment/home.index`

- 1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80
- 1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. "accepts" connection, notifying client
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`
3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket






# Nonpersistent HTTP (continued)

- 
- time
4. HTTP server closes TCP connection.
  5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
  6. Steps 1-5 repeated for each of 10 jpeg objects

# Nonpersistent HTTP (continued)



# Nonpersistent HTTP: Response Time


## Definition of Round-Trip Time

(RTT): time for a small packet to travel from client to server and back.

## Response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

total =  $2RTT + \text{transmit time}$






# Persistent HTTP

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection (for each connection, TCP buffers and TCP variables have to be allocated in the client and in the server)
- browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- Pipelining version:
  - client sends requests as soon as it encounters a referenced object
  - as little as one RTT for all the referenced objects



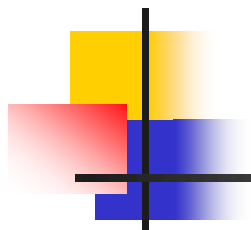
# Persistent HTTP

## Two versions of persistent connections

### ■ Without pipelining

- The client issues a new request only when the previous response has been received
- The client experiences one RTT in order to request and receive each of the referenced objects
- After the server sends an object over the persistent TCP connection, the connection idles while it waits for another request



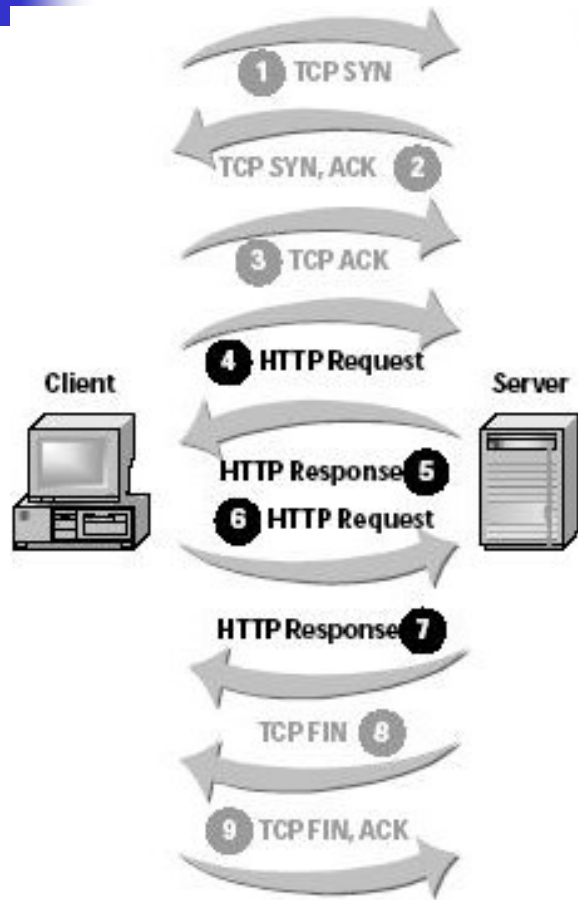


# Persistent HTTP

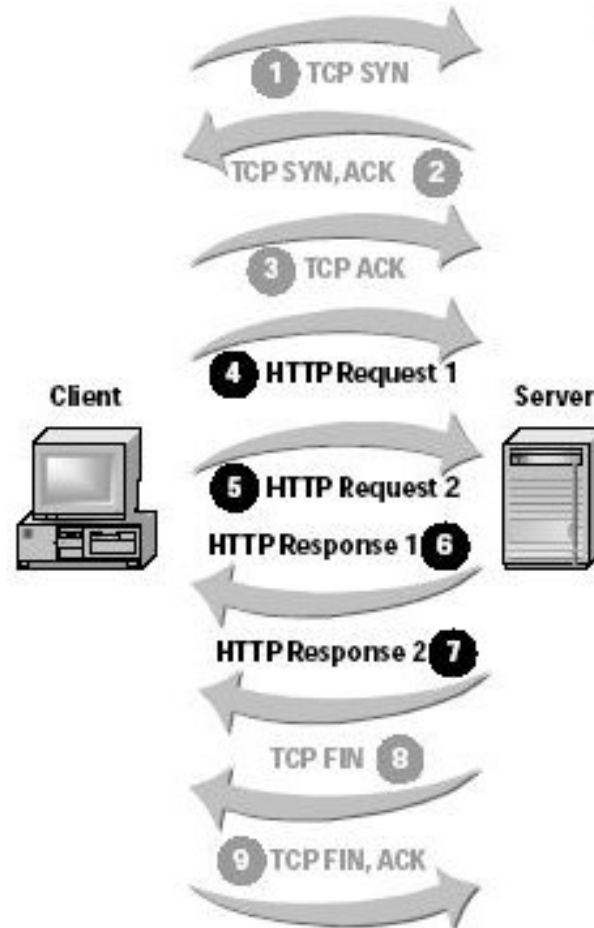
- With pipelining (default mode in HTTP/1.1)
  - The client issues a new request as soon as it encounters a reference.
  - It is possible for only one RTT to be expended for all the referenced objects
  - The pipelined TCP connection remains idle for a smaller fraction of time



# Persistent HTTP



Without Pipelining



With Pipelining



# HTTP Request Message

- two types of HTTP messages: *request, response*
- HTTP request message:
  - ASCII (human-readable format)

request line

(HTTP commands)

GET /somedir/page.html HTTP/1.1

header  
lines

Host: www.someschool.edu

User-agent: Mozilla/4.0

Connection: close


Accept-language: fr

Carriage return  
line feed

(extra carriage return, line feed)

indicates end  
of message





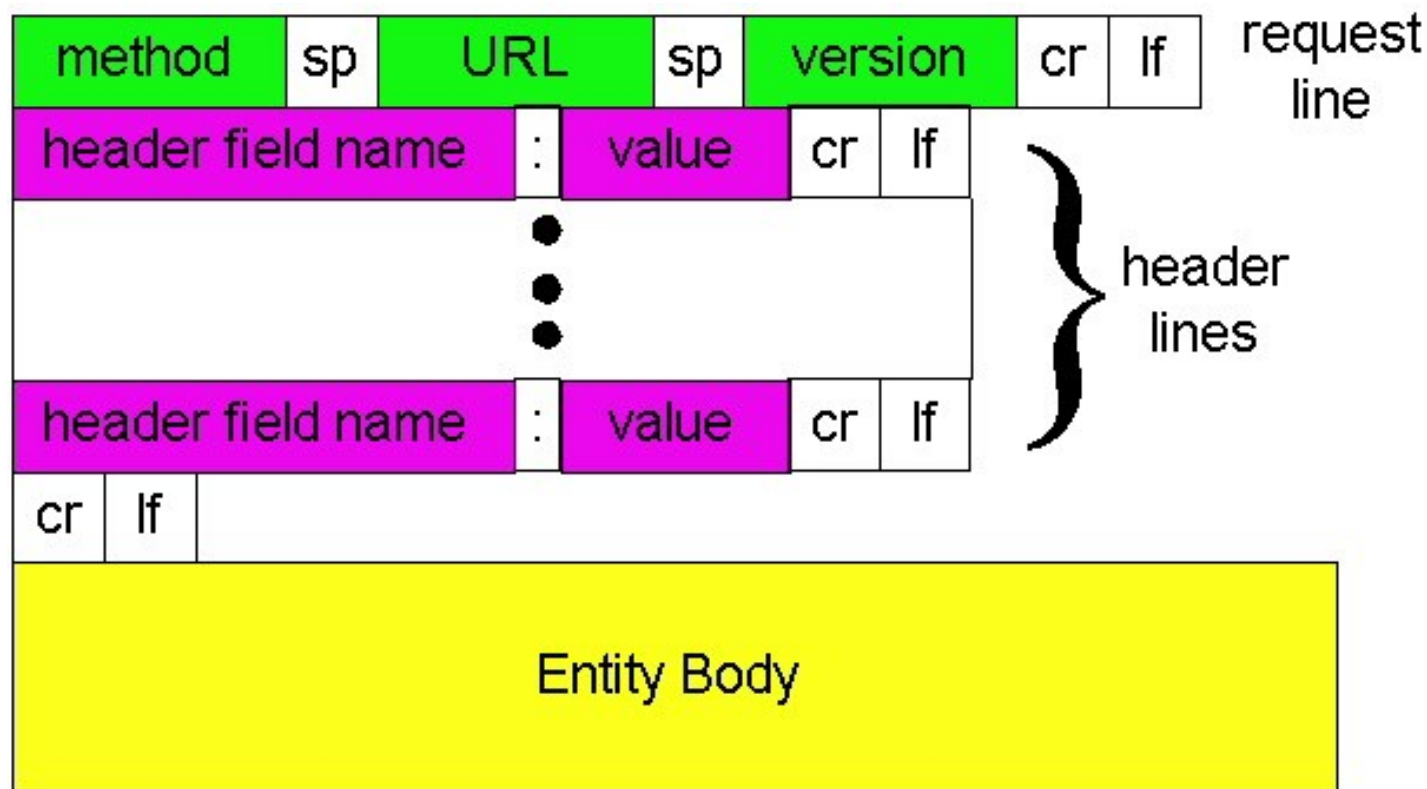
# HTTP Request Message

- Connection: close
  - The browser is telling the server that it does not want to use persistent connections (when it wants, it uses Keep-Alive)
- User-agent
  - Specifies the user agent, that is, the browser that is making the request to the server
  - Useful: the server could send different versions of the same object to different types of user agent
- Accept-language
  - The user prefers to receive a French version of the object, if such an object exists on the server; otherwise the server should send its default version



# HTTP Request Message

## General Format





# HTTP Request Message

---

- Web page often include form input

## Post method:

- Input is uploaded to server in entity body

## GET method:

- Input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`



# HTTP Request Message

## HTTP/1.0

- GET
- POST
- HEAD
  - asks server to leave requested object out of response


## HTTP/1.1

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field
- TRACE
  - invoke a remote, application-layer loop- back of the request message

# HTTP Request Message

## ■ GET


- Input in the URL field. No entity body
- Sure and idempotent (the side-effects of  $N > 0$  identical requests is the same as for a single request)
- Only ASCII characters
- Browser and proxy can cache the responses of the server



# HTTP Request Message

## ■ POST


- The URI specifies the script which should process the information
- The entity body contains what the user entered into the form fields
- The script is called at each form submission (no cache)



# HTTP Request Message

## ■ HEAD

- Similar to the GET method
- When a server receives a request with the HEAD method, it responds with an HTTP message but it leaves out the requested object
- Often used for debugging






# HTTP Request Message

## ■ PUT

- Used in conjunction with Web publishing tools
- Allows a user to upload an object to a specific path (directory) on a specific Web server (typically the object is a file)
- If the Request-URI refers to an already existing resource, the enclosed entity **SHOULD** be considered as a modified version of the one residing on the origin server.
- Used by applications which need to upload objects to Web servers



# HTTP Request Message

- DELETE


- Allows a user, or an application, to delete an object on a Web server



# HTTP Request Message


## ■ TRACE

- The TRACE method is used to invoke a remote, application-layer loop- back of the request message.
- The final recipient (origin server or proxy) of the request SHOULD reflect the message received back to the client as the entity-body of a 200 (OK) response.
- TRACE allows the client to see what is being received at the other end of the request chain and use that data for testing or diagnostic information.




# HTTP Request Message

- TRACE (example)



# HTTP Request Message

## ■ TRACE (example)





# HTTP Request Message

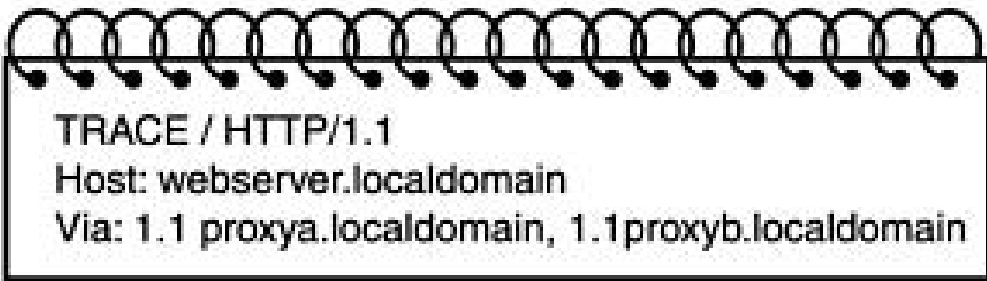
## ■ TRACE (example)

②



TRACE / HTTP/1.1  
Host: webserver.localdomain  
Via: 1.1 proxya.localdomain

③



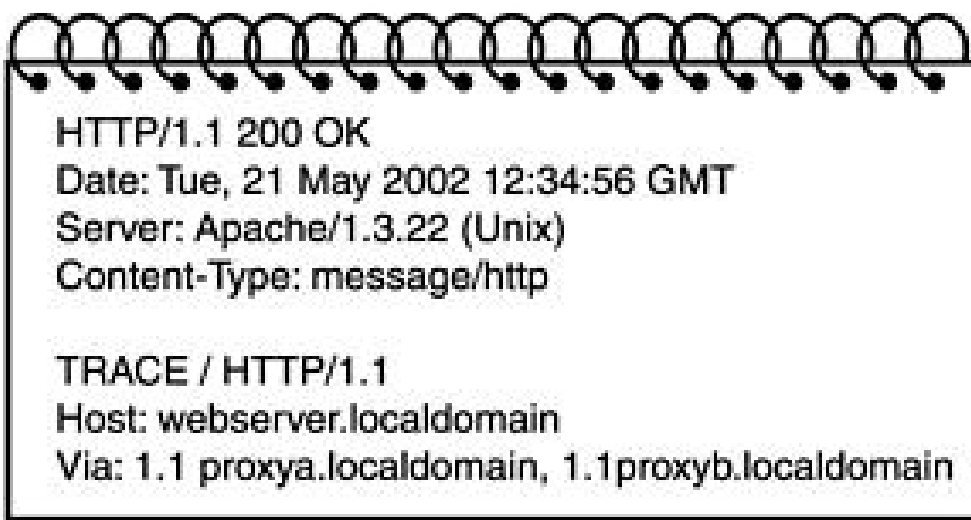
TRACE / HTTP/1.1  
Host: webserver.localdomain  
Via: 1.1 proxya.localdomain, 1.1 proxyb.localdomain



# HTTP Request Message

- TRACE (example)

④



```
HTTP/1.1 200 OK
Date: Tue, 21 May 2002 12:34:56 GMT
Server: Apache/1.3.22 (Unix)
Content-Type: message/http


TRACE / HTTP/1.1
Host: webserver.localdomain
Via: 1.1 proxya.localdomain, 1.1proxyb.localdomain
```



# HTTP Request Message

## ■ TRACE (example)


⑤



```
HTTP/1.1 200 OK
Date: Tue, 21 May 2002 12:34:56 GMT
Server: Apache/1.3.22 (Unix)
Content-Type: message/http
Via:1.1 proxyb.localdomain

TRACE / HTTP/1.1
Host: webserver.localdomain
Via: 1.1 proxya.localdomain, 1.1proxyb.localdomain
```





# HTTP Request Message

## ■ TRACE (example)

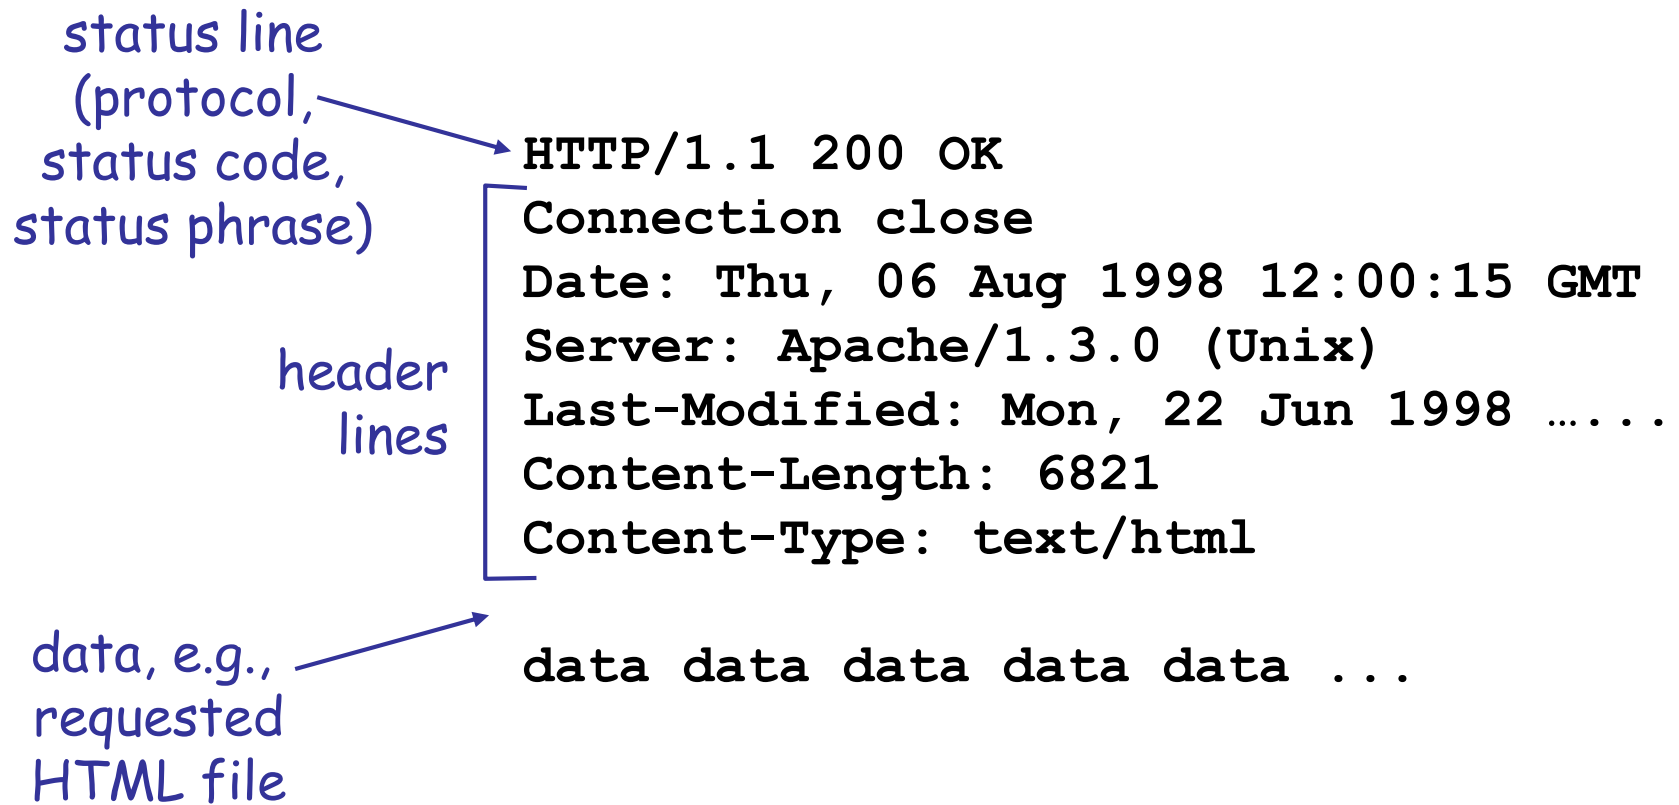
⑥

```
HTTP/1.1 200 OK
Date: Tue, 21 May 2002 12:34:56 GMT
Server: Apache/1.3.22 (Unix)
Content-Type: message/http
Via:1.1 proxyb.localdomain, 1.1 proxya.localdomain

TRACE / HTTP/1.1
Host: webserver.localdomain
Via: 1.1 proxya.localdomain, 1.1 proxyb.localdomain
```



# HTTP Response Message





# HTTP Response Status Codes

In first line in server->client response message.

A few sample codes:

## 200 OK

- request succeeded, requested object later in this message

## 301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

## 400 Bad Request

- request message not understood by server

## 404 Not Found

- requested document not found on this server

## 505 HTTP Version Not Supported



# HTTP Response Message

## Header Lines

- **Date:** indicates the time and the date when the HTTP response was created and sent by the server
- **Content-Type:** indicates that the object in the entity body is HTML text (The object type is officially indicated by the Content-Type: header and not by the file extension)



# Trying HTTP

## 1. Telnet to your favorite Web server:

```
telnet localhost 80
```

Opens TCP connection to port 80  
(default HTTP server port) at localhost.  
Anything typed in sent  
to port 80 at localhost

## 2. Type in a GET HTTP request:

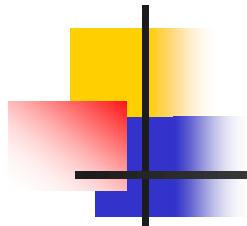
```
GET /form.html HTTP/1.1
```

```
Host: localhost
```

By typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

## 3. Look at response message sent by HTTP server!





# Trying HTTP (get.php)

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01//EN"  
"http://www.w3.org/TR/html4/strict.dtd">
```

```
<html lang="en">
```

```
<head>
```

```
<meta http-equiv="Content-type" content="text/html;  
charset=ISO-8859-1">
```

```
<title>send data using the GET method without forms and  
opening a new page</title>
```

```
</head>
```


```
<body>
```



# Trying HTTP

```
/* Redirect to a different page in the current directory that was
   requested */
<?php
    $host = $_SERVER['HTTP_HOST'];
    $uri  = rtrim(dirname($_SERVER['PHP_SELF']), '/\\');
    //rtrim — Strip whitespace (or other characters) from the
    end of a string
    $extra = 'receiveGET.php?name=Bob&age=21';
    header("Location: http://$host$uri/$extra");
    //Send a raw HTTP header
    exit;
?>
</body>
</html>
```





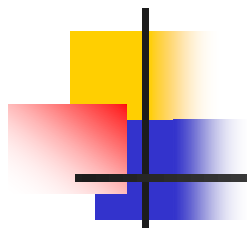
# Trying HTTP (receiveGet.php)

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01//EN"
"http://www.w3.org/TR/html4/strict.dtd">
<html>
  <head>
    <meta http-equiv="Content-Type" content="text/html;
    charset=iso-8859-1">
    <title>Receive data and forward to the client</title>
  </head>
  <body>
    <p> Hi <?php echo htmlspecialchars($_GET['name']); ?>,<br>
    You are <?php echo (int)$_GET['age']; ?> old.
  </p>
</body>
</html>
```

**htmlspecialchars — Convert special characters to HTML entities**







# Trying HTTP (post.php)

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01//EN"
"http://www.w3.org/TR/html4/strict.dtd">
<html lang="en">
  <head>
    <meta http-equiv="Content-type" content="text/html;
    charset=ISO-8859-1">
    <title>send data using the POST method without forms and
    opening a new page</title>
  </head>
    <body>
      <p>
```





# Trying HTTP (post.php)

---

```
<?php
$data =
sendPost('/receivePOST.php','name=Bob&age=21','localhost');
echo $data;

function sendPost($uri,$postdata,$host){
    $flow = fsockopen($host, 80, $errno, $errstr);
    //Initiates a socket connection to the resource specified by
    hostname.
    if (!$flow) {
        echo "$errstr ($errno)<br>\n";
        echo $flow;    }
    else {
```



# Trying HTTP (post.php)

```
$requestHTTP = "POST $uri HTTP/1.1\r\n";  
$requestHTTP.= "Host: $host\r\n";  
$requestHTTP.= "User-Agent: PHP Script\r\n";  
$requestHTTP.= "Content-Type: application/x-www-form-  
urlencoded\r\n";  
$requestHTTP.= "Content-Length: ".strlen($postdata)."\r\n";  
$requestHTTP.= "Connection: close\r\n\r\n";  
$requestHTTP.= $postdata;  
$replyHTTP = NULL;  
fwrite($flow, $requestHTTP);  
while (!feof($flow))  
    $replyHTTP.= fgets($flow, 128);  
$replyHTTP = explode("\r\n\r\n", $replyHTTP);  
$headerReplyHTTP = $replyHTTP[0];  
$contentsReplyHTTP = $replyHTTP[1];
```






# Trying HTTP (post.php)

```
if(!(strpos($headerReplyHTTP,"Transfer-Encoding:
    chunked")==false))
{ $temp=explode("\r\n",$contentsReplyHTTP);
  for($i=0;$i<count($temp);$i++)
    if($i==0 || ($i%2==0))
      $temp[$i]="";
  $contentsReplyHTTP=implode("",$temp);
  //implode — Join array elements with a string
}
fclose($flow);
return rtrim($contentsReplyHTTP);
}
} ?>
</p>
</body>
</html>
```





# Trying HTTP (receivePost.php)

Example of chunked transfer-encoding

HTTP/1.1 200 OK

Content-Type: text/plain

Transfer-Encoding: chunked

25\r\n

This is the data in the first chunk\r\n

1C\r\n

and this is the second one\r\n

3\r\n

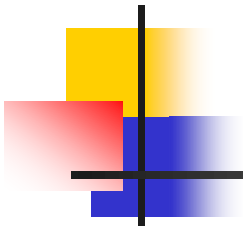
Con\r\n

8\r\n

Sequence\r\n

0\r\n





# Trying HTTP (receivePost.php)

Hi <?php echo htmlspecialchars(\$\_POST['name']); ?>,  
<br> You are <?php echo (int)\$\_POST['age']; ?> old

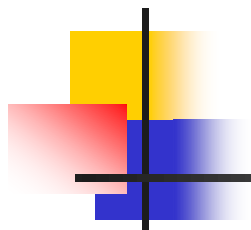




# Header Lines

---

- How does a browser decide which header lines to include in a request message?
  - Browser type and version
  - User configuration of the browser
  - If the browser has a cached, but possibly out-of-date version of the object
- How does a Web server decide which header lines to include in a response message?
  - Server type and version
  - Server configuration



# Header Lines


- HTTP/1.0 defines 16 different header lines (no mandatory)
- HTTP/1.1 defines 51 different header lines (only one Host is mandatory) in the request messages to avoid that the server replies 400 Bad Request.





# Authorization and Cookies

- HTTP server is stateless
  - Simplifies server design
  - High performance Web servers that can handle thousands of simultaneous TCP connections
- Sometimes it is desirable for a Web site to identify users
  - To restrict user access
  - To serve a content as a function of the user identity



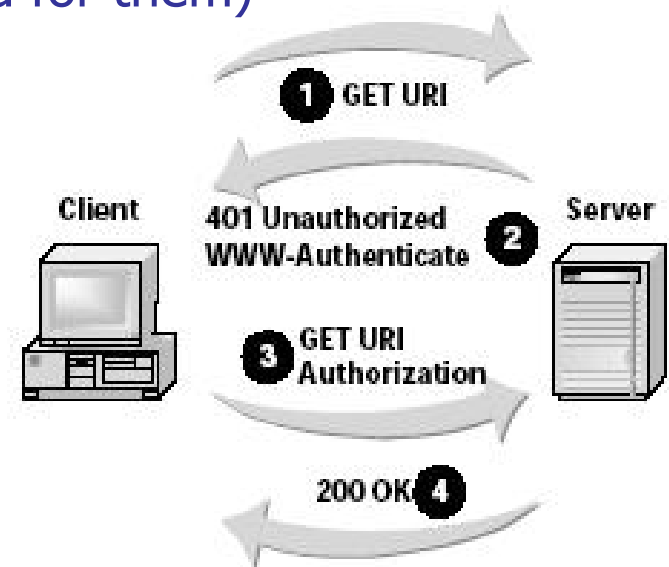
# Authorization

- Requesting and receiving authorizations is often done by using special HTTP headers and status codes
- Scenario
  1. Client requests an object from a server and the server requires an authorization
  2. Client sends an ordinary request message with no special header lines
  3. The server responds with an empty entity body and 401 Authorization Required status code. In the response message the server includes the WWW-Authenticate: header

# Authorization

4. The client prompts the user for a username and passwd
5. The client resends the request message including an Authorization: header line
6. After obtaining the first object, the client continues to send username and password in subsequent requests for objects on the server (username and passwords are cached and therefore the user is not prompted for them)

- Note: this is a very weak form of authorization
  - One can sniff (read and store) all the packets and therefore steal the login password.





# Cookies

Many major Web sites use cookies

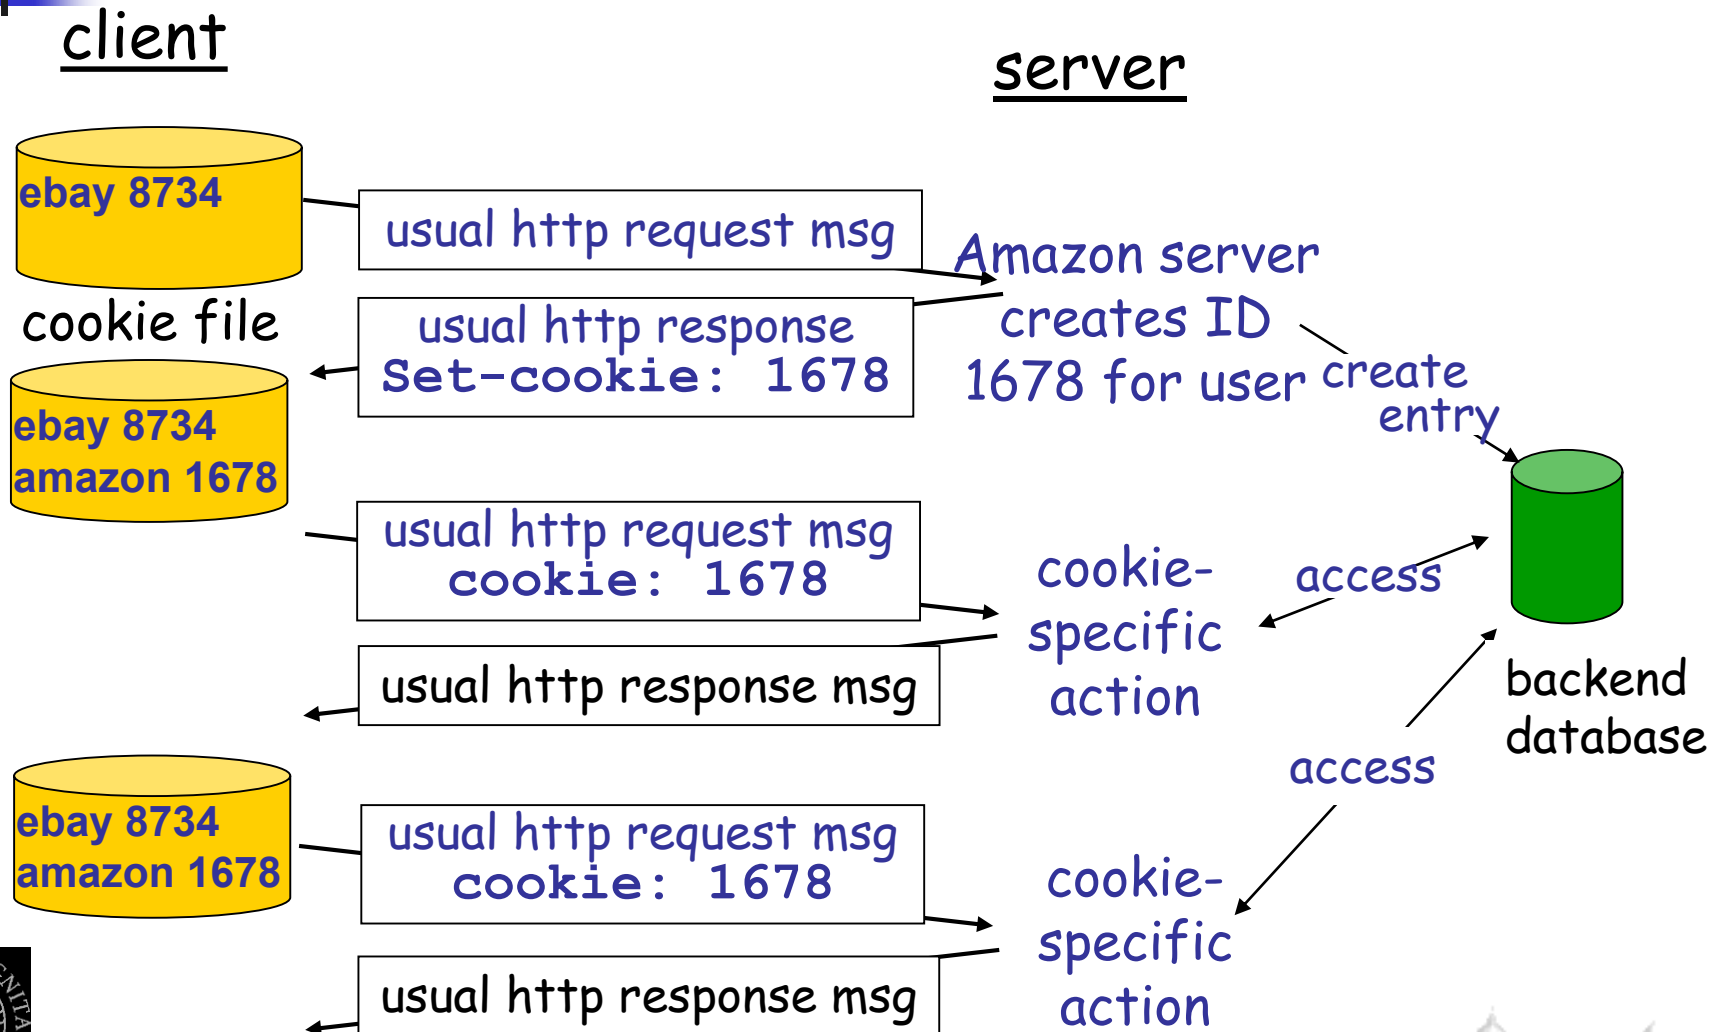
## Four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

## Example:

- Susan always access Internet always from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

# Cookies



# Cookies





# Cookies

aside

## What cookies can bring:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

## Cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites


## How to keep "state":

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

# Web caching


**Goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - object in cache: cache returns object
  - else cache requests object from origin server, then returns object to client





# More about Web caching





# More about Web caching

- cache acts as both client and server
- typically cache is installed by the Internet Service Providers (university, company, residential ISP)

## Why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link.
- Internet dense with caches: enables "poor" content providers to effectively deliver content (but so does P2P file sharing)


# Caching Example


## Assumptions

- average object size = 1,000,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router (Internet delay) = 2 sec

## Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + milliseconds





# Caching Example

## Some Computation

Lan Traffic Intensity

$$(15 \text{ requests/sec}) * (1\text{Mb/request}) / 100\text{Mbps} = 0.15$$

An intensity of 0.15 results in, at most, tens of milliseconds of delay

Access Link Intensity

$$(15 \text{ requests/sec}) * (1\text{Mb/request}) / 15\text{Mbps} = 1$$

An intensity of 1 results in a very large delay of the order of minutes




# Caching Example (continued)

## possible solution

- increase bandwidth of access link to, say, 100 Mbps

## consequence

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay  
= 2 sec + msecs + msecs
- often a costly upgrade



# Caching Example (continued)


## possible solution: install cache

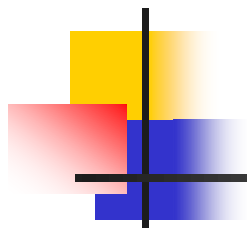
- suppose hit rate is 0.4

## consequence

- 40% requests will be satisfied almost immediately (order of milliseconds)
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + access delay + LAN delay  

$$= .6 \cdot (2.01) \text{ secs} + .4 \cdot \text{milliseconds} < 1.4 \text{ secs}$$





# Problem

- Web caches can reside in a client (managed by the user's browser) or in an intermediate network cache server
- Problem: a copy of an object residing in the cache may be *stale* (object housed in the Web server may have been modified since the copy was cached at the client)
- Solution: the conditional GET
  - The request message uses the GET method
  - The request message includes an If-Modified-Since: header line



# Conditional Cache

- **Goal:** don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request  
**If-modified-since:**  
**<date>**
- server: response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**

cache

server

